

VEERESH KOLIWAD

602-451-5421 | veereshkoliwad99@gmail.com | [linkedin.com/in/veereshkoliwad](https://www.linkedin.com/in/veereshkoliwad) | github.com/veereshgit99

SUMMARY

Software Engineer with 2.5+ years of experience building distributed backend systems and deploying scalable cloud solutions. Led ML-driven automation, improved API performance by 30%, and built microservices that cut order processing time by 90%. Passionate about system design, performance tuning, and real-world ML applications.

EXPERIENCE

Research Assistant

Jan 2025 – Present

Arizona State University

Tempe, AZ

- Evaluated MiniLLM on benchmark datasets (OpenOrca, AlpacaEval) to explore trade-offs in LLM compression through knowledge distillation.
- Reduced inference latency by 45% using LayerDrop and pruning techniques, making compressed LLMs more practical for deployment.

Associate Software Developer

Jul 2021 – Dec 2023

SAP LABS

Bangalore, India

- Built a purchase order microservice for SAP ERP to automate manual workflows, reducing processing time by 90% for 5K+ monthly orders.
- Migrated legacy applications to SAP Cloud, cutting infrastructure costs by 20% and improving scalability.
- Optimized API response time by 30% using Redis caching, improving overall user engagement by 20%.
- Led development of a bug prediction tool with a 4-member team, reducing support tickets by 25%.

Software Development Intern

Feb 2021 – Jul 2021

SAP LABS

Bangalore, India

- Developed a Jira workflow gadget using JavaScript and HTML, used by 10+ teams to simplify ticket tracking.
- Built and deployed full-stack features using React and MongoDB on AWS, for SAP's Bydesign module.

PROJECTS

Football Analysis System | Python, YOLOv8, OpenCV, KMeans, Optical Flow

Jun 2024 – Aug 2024

- Built a real-time player/ball detection system using YOLOv8, achieving 95% accuracy on Bundesliga datasets.
- Implemented KMeans clustering for team identification and optical flow for tracking player movements.
- Applied perspective transformation to standardize camera angles across match footage.

RAG Chatbot | LangChain, Hugging Face, ChromaDB, Vector Search

Jan 2025

- Developed a context-aware chatbot using Retrieval-Augmented Generation (RAG) to integrate real-time data with LLMs.
- Optimized response accuracy by implementing vector search and data chunking in ChromaDB.
- Deployed the system with Flask API for seamless integration with frontend applications.

EDUCATION

Arizona State University

Jan 2024 – Dec 2025

Master's in Computer Science

Tempe, AZ

RV College of Engineering

Aug 2017 – May 2021

Bachelor's in Computer Science

Bangalore, India

TECHNICAL SKILLS

Languages: Python, C/C++, Java (Intermediate), SQL, JavaScript, HTML/CSS

Machine Learning: NLP, Computer Vision, LLMs, Recommendation Systems

Tools: Docker, Kubernetes, React, Kafka, Elastic Search, Redis, Git

Cloud & Databases: AWS, SAP Cloud Platform, PostgreSQL, MongoDB

Methodologies: REST APIs, Microservices, Distributed Systems, Agile, CI/CD